

Code for slide deck on penalised regression and cross-validation

Benjamin Skov Kaas-Hansen

11/02/2020

Contents

Setup	2
Lasso regression example: biopsies from breast cancer patients	2
Lasso regression	2
Ridge and elastic net models	5
Over-fitting biopsy	8
Delassoing	10
Cross-validation Pima	12
Homegrown	12
Using packages	12
Cross-validation biopsy	12
Test out-of-sample performance of the CV model	13
Resampling methods: jackknife and bootstrap	13
Sampling distribution	13
Jackknife	13
Exercise: cross-validation	13
1. LOO-CV error rate	13
2. Use proper cost function	13
3. Difference between error rates, and their interpretation	13
4. 10-fold CV	13

Exercise: penalised regression	13
1. + 2. Lasso regression	13
3. Why does it normally make sense to normalise predictors	13
4. Using CV to get reasonable estimate of λ	13
5. Obtain coefficients for “best” λ	13
6. Re-fit with correct family	13
7. As ridge regression	13
8. Get idea about sparse solution using ridge results?	13
9. Elastic net ($\alpha = 0.5$)	13
10. Delasso the results	13
11. Selective inference	13

Setup

```
packages <- c("plyr", "tidyr", "glmnet", "selectiveInference", "MASS", "tidyverse")
for (p in packages)
  library(p, character.only = TRUE)
knitr::opts_chunk$set(fig.align = "center")

theme_set(theme_minimal())

# Little helper to get the glmnet coefficients in nice tidy format
pretty_coefs <- function(coefs) { # coefs: the output from coef(fit_object, s = [value])
  enframe(coefs[, 1], "predictor", "coefficient") %>%
    filter(coefficient != 0) %>%
    arrange(desc(abs(coefficient)))
}
```

Lasso regression example: biopsies from breast cancer patients

Lasso regression

Let's look at the data


```
data(biopsy)
summary(biopsy) # NA's in V6; mean varies across variables but anyway somewhere around 2 and 4
```

```
##      ID          V1          V2          V3
## Length:699      Min.   : 1.000   Min.   : 1.000   Min.   : 1.000
## Class :character 1st Qu.: 2.000   1st Qu.: 1.000   1st Qu.: 1.000
## Mode  :character Median : 4.000   Median : 1.000   Median : 1.000
##              Mean   : 4.418   Mean   : 3.134   Mean   : 3.207
##              3rd Qu.: 6.000   3rd Qu.: 5.000   3rd Qu.: 5.000
##              Max.   :10.000   Max.   :10.000   Max.   :10.000
##
```

```
##           V4           V5           V6           V7
## Min.      : 1.000    Min.      : 1.000    Min.      : 1.000    Min.      : 1.000
## 1st Qu.: 1.000    1st Qu.: 2.000    1st Qu.: 1.000    1st Qu.: 2.000
## Median : 1.000    Median : 2.000    Median : 1.000    Median : 3.000
## Mean      : 2.807    Mean      : 3.216    Mean      : 3.545    Mean      : 3.438
## 3rd Qu.: 4.000    3rd Qu.: 4.000    3rd Qu.: 6.000    3rd Qu.: 5.000
## Max.      :10.000    Max.      :10.000    Max.      :10.000    Max.      :10.000
##
##                               NA's      :16
##           V8           V9           class
## Min.      : 1.000    Min.      : 1.000    benign      :458
## 1st Qu.: 1.000    1st Qu.: 1.000    malignant:241
## Median : 1.000    Median : 1.000
## Mean      : 2.867    Mean      : 1.589
## 3rd Qu.: 4.000    3rd Qu.: 1.000
## Max.      :10.000    Max.      :10.000
##
```


```
biopsy_complete <- na.exclude(biopsy) # remove rows with any missing value
biopsy_predictors <- select(biopsy_complete, -ID, -class) %>%
  scale() # note attributes "remember" normalisation factors; useful for transforming test set

bind_rows(gather(as_tibble(biopsy_predictors), var, value) %>%
  mutate(scale = "normalised"),
  gather(select(biopsy_complete, -ID, -class), var, value) %>%
  mutate(scale = "original")) %>%
ggplot(aes(x = value, colour = scale)) +
  geom_density(position = "identity") +
  scale_x_continuous(breaks = -2:10) +
  facet_wrap(~ var, scales = "free_y") +
  theme(axis.text.y = element_blank())
```



```
lasso_logreg <- glmnet(biopsy_predictors, biopsy_complete$class, family = "binomial")

# Coefficient profile plot (built-in: ugly but easy)
plot(lasso_logreg, xvar = "lambda", label = TRUE, lwd = 1.5)
```




```
pretty_coefs(coef(lasso_logreg, s = exp(-1.2)))
```

```
## # A tibble: 4 x 2
##   predictor coefficient
##   <chr>         <dbl>
## 1 (Intercept)  -0.642
## 2 V6           0.219
## 3 V2           0.138
## 4 V3           0.109
```

BACK TO PRESENTATION

Ridge and elastic net models


```
ridge_logreg <- update(lasso_logreg, alpha = 0)
plot(ridge_logreg, xvar = "lambda")
```



```
pretty_coefs(coef(ridge_logreg, s = exp(2))) # all shrunk but no real ranking or anything
```

```
## # A tibble: 10 x 2
##   predictor coefficient
##   <chr>          <dbl>
## 1 (Intercept)    -0.630
## 2 V6              0.0457
## 3 V3              0.0448
## 4 V2              0.0447
## 5 V7              0.0414
## 6 V1              0.0395
## 7 V8              0.0391
## 8 V4              0.0383
## 9 V5              0.0373
## 10 V9             0.0224
```

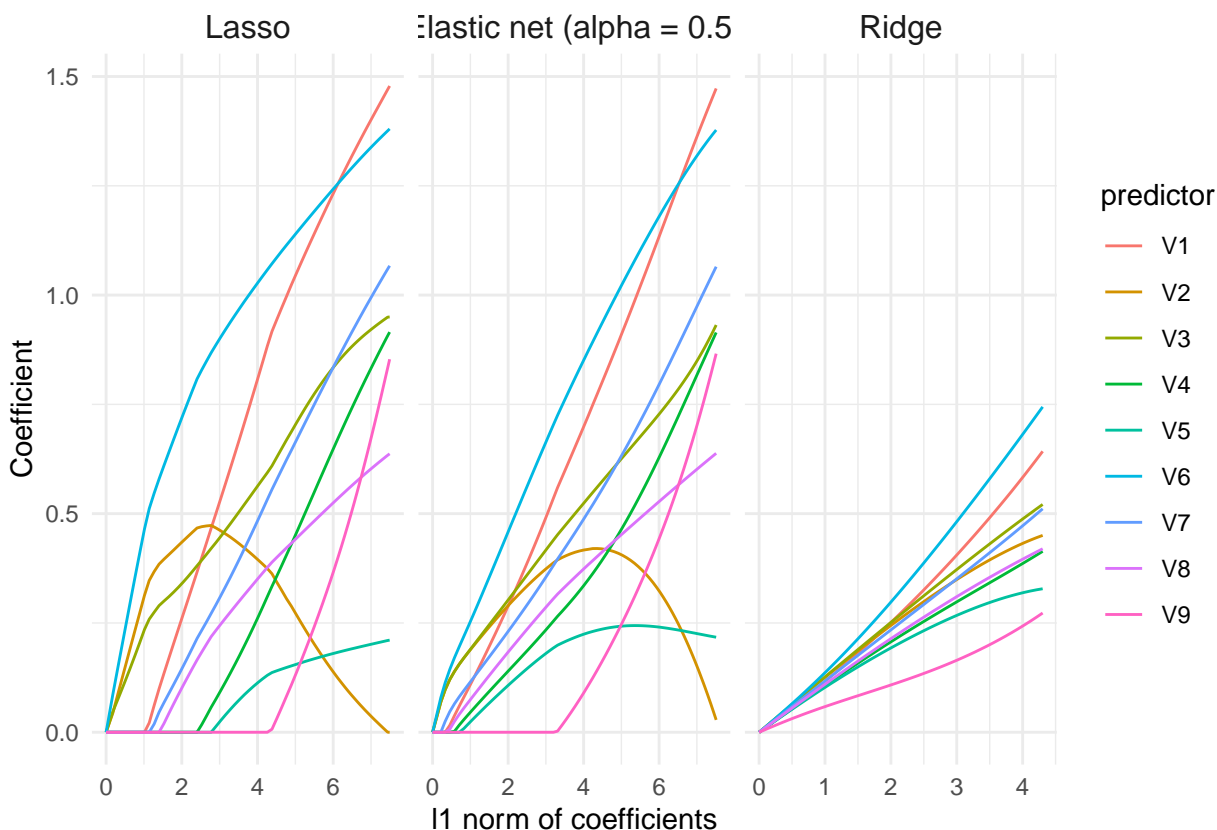
```
elastic_logreg <- update(lasso_logreg, alpha = 0.5)
plot(elastic_logreg, xvar = "lambda")
```



```
pretty_coefs(coef(elastic_logreg, s = exp(-1.2))) # all shrunk AND "ranking"
```

```
## # A tibble: 9 x 2
##   predictor coefficient
##   <chr>          <dbl>
## 1 (Intercept)   -0.719
## 2 V6            0.298
## 3 V3            0.215
## 4 V2            0.213
## 5 V1            0.148
## 6 V7            0.141
## 7 V8            0.0996
## 8 V4            0.0675
## 9 V5            0.0435
```

```
ldply(list(Lasso = lasso_logreg, Ridge = ridge_logreg, Elastic_net = elastic_logreg),
      function(.) with(., data.frame(as.matrix(t(beta)), x = apply(abs(beta), 2, sum))), .id = "mod") %>%
  gather(predictor, value, -x, -mod) %>%
  mutate(mod = factor(mod, levels = c("Lasso", "Elastic_net", "Ridge"),
    labels = c("Lasso", "Elastic net (alpha = 0.5)", "Ridge"))) %>%
  ggplot(aes(x, value, colour = predictor)) +
  geom_line() +
  labs(x = "l1 norm of coefficients", y = "Coefficient") +
  facet_wrap(~ mod, scales = "free_x") +
  theme(strip.text = element_text(size = 12))
```



Over-fitting biopsy

It's a little data set but let's try and split into training and test sets. There is quite some over-fitting on the relative scale, but the absolute prediction error difference is minor.

```
set.seed(42) # reproducible stochastic code
train_idx <- runif(nrow(biopsy_complete)) <= 0.8 # not exactly indices
mean(train_idx) # fraction in training set

## [1] 0.806735

biopsy_train <- filter(biopsy_complete, train_idx)
biopsy_test <- filter(biopsy_complete, !train_idx)
all_equal(biopsy_complete, bind_rows(biopsy_train, biopsy_test)) # sanity check

## [1] TRUE

# Alternative with actual indices (mostly a matter of taste)
set.seed(42)
train_idx <- which(runif(nrow(biopsy_complete)) <= 0.8)
length(train_idx) / nrow(biopsy_complete)

## [1] 0.806735
```



```

biopsy_train <- slice(biopsy_complete, train_idx)
biopsy_test <- slice(biopsy_complete, -train_idx)
all_equal(biopsy_complete, bind_rows(biopsy_train, biopsy_test)) # sanity check

## [1] TRUE

# Normalise predictors and put them in matrix format
predictors_train <- select(biopsy_train, -ID, -class) %>%
  scale()
predictors_test <- select(biopsy_test, -ID, -class) %>%
  scale()

# Train model
lasso_biopsy_train <- glmnet(predictors_train, biopsy_train$class, family = "binomial")
D_train <- predict(lasso_biopsy_train, predictors_train, type = "class") %>%
  apply(2, function(.) mean(. != biopsy_train$class)) # s39 is the best-performing


D_test <- predict(lasso_biopsy_train, predictors_test, type = "class") %>%
  apply(2, function(.) mean(. != biopsy_test$class))

tibble(D_test = D_test[which.min(D_train)],
       D_train = D_train[which.min(D_train)],
       D_diff_abs = scales::percent(D_test - D_train),
       D_diff_rel = scales::percent((D_test - D_train) / D_train, big.mark = ","))

## # A tibble: 1 x 4
##   D_test D_train D_diff_abs D_diff_rel
##   <dbl>   <dbl> <chr>      <chr>
## 1 0.0379 0.0272 1%          39%


ggplot(mapping = aes(x = log(lasso_biopsy_train$lambda))) +
  coord_cartesian(ylim = c(0, NA)) +
  geom_line(aes(y = D_train, colour = "Within-sample")) +
  geom_line(aes(y = D_test, colour = "Hold-out")) +
  labs(y = "log(prediction error)", x = expression(log(lambda))) +
  scale_x_reverse()

```



Delassoing

```
plot(lasso_logreg, xvar = "lambda")
```



```
lambda <- exp(-4.5)
beta <- coef(lasso_logreg, x = biopsy_predictors, y = biopsy_complete$class,
             s = lambda/nrow(biopsy_complete), exact = TRUE)
delasso_fit <- fixedLassoInf(biopsy_predictors, as.numeric(biopsy_complete$class)-1, beta, lambda, "binomial",
                           alpha = 0.05)
```

```
## Warning in fixedLogitLassoInf(x, y, beta, lambda, alpha = alpha, type =
## type, : Solution beta does not satisfy the KKT conditions (to within specified
## tolerances)
```

```
delasso_fit
```

```
##
## Call:
## fixedLassoInf(x = biopsy_predictors, y = as.numeric(biopsy_complete$class) -
## 1, beta = beta, lambda = lambda, family = "binomial", alpha = 0.05)
##
## Testing results at lambda = 0.011, with alpha = 0.050
##
## Var    Coef Z-score P-value LowConfPt UpConfPt LowTailArea UpTailArea
## 1  1.509  3.770  0.596   -9.749   2.005    0.025    0.024
## 2 -0.019 -0.030  0.989    0.857    Inf     0.025    0.000
## 3  0.964  1.399  0.962   -Inf     0.570    0.000    0.025
## 4  0.947  2.680  0.808  -17.165   1.237    0.025    0.024
## 5  0.215  0.617  0.863  -10.079   0.635    0.025    0.024
## 6  1.396  4.084  0.000    0.743   3.114    0.025    0.025
## 7  1.095  2.611  0.749  -14.329   1.537    0.025    0.025
## 8  0.650  1.889  0.805  -12.461   1.008    0.025    0.025
## 9  0.927  1.629  0.705  -10.952   1.724    0.025    0.025
```

```
##  
## Note: coefficients shown are full regression coefficients  
  
knitr::opts_chunk$set(include = FALSE)
```

Cross-validation Pima

Homegrown

Using packages

The `modelr`-package has some powerful functionalities for CV, LOOCV and bootstrapping.

Cross-validation biopsy

Use cross-validation to find best λ value. We found quite clear over-fitting above. Let's try to remedy this with cross-validation.

Test out-of-sample performance of the CV model

Resampling methods: jackknife and bootstrap

Sampling distribution

Jackknife

Exercise: cross-validation

1. LOO-CV error rate
2. Use proper cost function
3. Difference between error rates, and their interpretation
4. 10-fold CV

Exercise: penalised regression

1. + 2. Lasso regression
3. Why does it normally make sense to normalise predictors
4. Using CV to get reasonable estimate of λ
5. Obtain coefficients for “best” λ
6. Re-fit with correct family

Quite some difference between the sets of predictors kept:

7. As ridge regression
8. Get idea about sparse solution using ridge results?
9. Elastic net ($\alpha = 0.5$)

Ideally, one should do CV over α as well.

10. Delasso the results
11. Selective inference