

Code for slide deck on penalised regression and cross-validation

Benjamin Skov Kaas-Hansen

11/02/2020

Contents

Setup	2
Cross-validation Pima	2
Homegrown	2
Using packages	3
Lasso regression example: biopsies from breast cancer patients	4
Lasso regression	4
Ridge and elastic net models	6
Over-fitting biopsy	9
Delassoing	11
Cross-validation	13
Test out-of-sample performance of the CV model	14
Cross-validation to pick best combination of λ and α	14
Resampling methods: jackknife and bootstrap	15
Sampling distribution	15
Jackknife	18
Exercise: cross-validation	19
1. LOO-CV error rate	20
2. Use proper cost function	20
3. Difference between error rates, and their interpretation	20
4. 10-fold CV	21

Exercise: penalised regression	22
1. + 2. Lasso regression	22
3. Why does it normally make sense to normalise predictors	22
4. Using CV to get reasonable estimate of λ	22
5. Obtain coefficients for “best” λ	22
6. Re-fit with correct family	22
7. As ridge regression	22
8. Get idea about sparse solution using ridge results?	22
9. Elastic net ($\alpha = 0.5$)	22
10. Delasso the results	22
11. Selective inference	22

Setup

```
packages <- c("plyr", "tidyr", "broom", "boot", "glmnet", "selectiveInference", "MASS", "tidyverse")
for (p in packages)
  library(p, character.only = TRUE)
knitr::opts_chunk$set(fig.align = "center")

theme_set(theme_minimal())

# Little helper to get the glmnet coefficients in nice tidy format
pretty_coefs <- function(coefs) { # coefs: the output from coef(fit_object, s = [value])
  enframe(coefs[, 1], "predictor", "coefficient") %>%
    filter(coefficient != 0) %>%
    arrange(desc(abs(coefficient)))
}
```

Cross-validation Pima

Homegrown

```
# Very simplistic implementation
data(PimaIndiansDiabetes2, package = "mlbench")
glimpse(PimaIndiansDiabetes2)

## Rows: 768
## Columns: 9
## $ pregnant <dbl> 6, 1, 8, 1, 0, 5, 3, 10, 2, 8, 4, 10, 10, 1, 5, 7, 0, 7, 1...
## $ glucose <dbl> 148, 85, 183, 89, 137, 116, 78, 115, 197, 125, 110, 168, 1...
## $ pressure <dbl> 72, 66, 64, 66, 40, 74, 50, NA, 70, 96, 92, 74, 80, 60, 72...
## $ triceps <dbl> 35, 29, NA, 23, 35, NA, 32, NA, 45, NA, NA, NA, NA, 23, 19...
## $ insulin <dbl> NA, NA, NA, 94, 168, NA, 88, NA, 543, NA, NA, NA, NA, 846,...
## $ mass <dbl> 33.6, 26.6, 23.3, 28.1, 43.1, 25.6, 31.0, 35.3, 30.5, NA, ...
## $ pedigree <dbl> 0.627, 0.351, 0.672, 0.167, 2.288, 0.201, 0.248, 0.134, 0....
## $ age <dbl> 50, 31, 32, 21, 33, 30, 26, 29, 53, 54, 30, 34, 57, 59, 51...
## $ diabetes <fct> pos, neg, pos, neg, pos, neg, pos, neg, pos, pos, neg, pos...
```

```
summary(PimaIndiansDiabetes2)
```

```
##      pregnant      glucose      pressure      triceps
## Min.   : 0.000   Min.   : 44.0   Min.   : 24.00   Min.   : 7.00
## 1st Qu.: 1.000   1st Qu.: 99.0   1st Qu.: 64.00   1st Qu.:22.00
## Median : 3.000   Median :117.0   Median : 72.00   Median :29.00
## Mean   : 3.845   Mean   :121.7   Mean   : 72.41   Mean   :29.15
## 3rd Qu.: 6.000   3rd Qu.:141.0   3rd Qu.: 80.00   3rd Qu.:36.00
## Max.   :17.000   Max.   :199.0   Max.   :122.00   Max.   :99.00
##
##      NA's      :5      NA's      :35      NA's      :227
##      insulin      mass      pedigree      age      diabetes
## Min.   : 14.00   Min.   :18.20   Min.   :0.0780   Min.   :21.00   neg:500
## 1st Qu.: 76.25   1st Qu.:27.50   1st Qu.:0.2437   1st Qu.:24.00   pos:268
## Median :125.00   Median :32.30   Median :0.3725   Median :29.00
## Mean   :155.55   Mean   :32.46   Mean   :0.4719   Mean   :33.24
## 3rd Qu.:190.00   3rd Qu.:36.60   3rd Qu.:0.6262   3rd Qu.:41.00
## Max.   :846.00   Max.   :67.10   Max.   :2.4200   Max.   :81.00
## NA's      :374      NA's      :11
```

```
set.seed(42)
```

```
pima <- na.exclude(PimaIndiansDiabetes2) %>%
  mutate(cv_fold = sample(1:10, n(), replace = TRUE))
table(pima$cv_fold) # fairly equal distribution
```

```
##
##  1  2  3  4  5  6  7  8  9 10
## 44 40 29 39 39 47 29 26 46 53
```

```
err_cv <- c()
for (i in unique(pima$cv_fold)) {
  train <- filter(pima, cv_fold != i)
  mod <- glm(diabetes ~ age + mass + insulin + pregnant, data = train, family = binomial)

  val <- filter(pima, cv_fold == i)
  y_pred <- predict(mod, newdata = val, type = "response")
  y_true <- as.numeric(val$diabetes) - 1 # bring binary factor to 0/1 scale
  err <- mean(abs(y_true - y_pred) > 0.5)
  err_cv <- c(err_cv, err)
}
err_cv
```

```
## [1] 0.2500000 0.2564103 0.2608696 0.2641509 0.3333333 0.2500000 0.3076923
## [8] 0.3103448 0.3448276 0.2978723
```

```
mean(err_cv)
```

```
## [1] 0.2875501
```

Using packages

```
binary_pred_cost <- function(y_true, y_pred) {
  mean(abs(y_true - y_pred) > 0.5)
```

```

}

library(boot)
pima_glm <- glm(diabetes ~ age + mass + insulin + pregnant, data = pima, family = binomial)

pima_loo <- cv.glm(pima, pima_glm, cost = binary_pred_cost)
pima_loo$delta # 1st is raw estimate, 2nd is bias-corrected

## [1] 0.2857143 0.2844908

pima_cv1 <- cv.glm(pima, pima_glm, cost = binary_pred_cost, K = 10)
pima_cv1$delta # 1st is raw estimate, 2nd is bias-corrected

## [1] 0.2806122 0.2801112

# This model doesn't really overfit

```

The modelr-package has some powerful functionalities for CV, LOOCV and bootstrapping.

Lasso regression example: biopsies from breast cancer patients

Lasso regression

Let's look at the data

```

data(biopsy)
summary(biopsy) # NA's in V6; mean varies across variables but anyway somewhere around 2 and 4

```


	ID	V1	V2	V3
##	Length:699	Min. : 1.000	Min. : 1.000	Min. : 1.000
##	Class :character	1st Qu.: 2.000	1st Qu.: 1.000	1st Qu.: 1.000
##	Mode :character	Median : 4.000	Median : 1.000	Median : 1.000
##		Mean : 4.418	Mean : 3.134	Mean : 3.207
##		3rd Qu.: 6.000	3rd Qu.: 5.000	3rd Qu.: 5.000
##		Max. :10.000	Max. :10.000	Max. :10.000
##				
##	V4	V5	V6	V7
##	Min. : 1.000	Min. : 1.000	Min. : 1.000	Min. : 1.000
##	1st Qu.: 1.000	1st Qu.: 2.000	1st Qu.: 1.000	1st Qu.: 2.000
##	Median : 1.000	Median : 2.000	Median : 1.000	Median : 3.000
##	Mean : 2.807	Mean : 3.216	Mean : 3.545	Mean : 3.438
##	3rd Qu.: 4.000	3rd Qu.: 4.000	3rd Qu.: 6.000	3rd Qu.: 5.000
##	Max. :10.000	Max. :10.000	Max. :10.000	Max. :10.000
##			NA's :16	
##	V8	V9	class	
##	Min. : 1.000	Min. : 1.000	benign :458	
##	1st Qu.: 1.000	1st Qu.: 1.000	malignant:241	
##	Median : 1.000	Median : 1.000		
##	Mean : 2.867	Mean : 1.589		
##	3rd Qu.: 4.000	3rd Qu.: 1.000		
##	Max. :10.000	Max. :10.000		
##				

```

biopsy_complete <- na.exclude(biopsy) # remove rows with any missing value
biopsy_predictors <- select(biopsy_complete, -ID, -class) %>%
  scale() # note attributes "remember" normlisation factors; useful for transforming test set

bind_rows(gather(as_tibble(biopsy_predictors), var, value) %>%
  mutate(scale = "normalised"),
  gather(select(biopsy_complete, -ID, -class), var, value) %>%
  mutate(scale = "original")) %>%
ggplot(aes(x = value, colour = scale)) +
  geom_density(position = "identity") +
  scale_x_continuous(breaks = -2:10) +
  facet_wrap(~ var, scales = "free_y") +
  theme(axis.text.y = element_blank())

```




```

lasso_logreg <- glmnet(biopsy_predictors, biopsy_complete$class, family = "binomial")

# Coefficient profile plot (built-in: ugly but easy)
plot(lasso_logreg, xvar = "lambda", label = TRUE, lwd = 1.5)

```




```
pretty_coefs(coef(lasso_logreg, s = exp(-1.2)))
```

```
## # A tibble: 4 x 2
##   predictor coefficient
##   <chr>         <dbl>
## 1 (Intercept)  -0.642
## 2 V6           0.219
## 3 V2           0.138
## 4 V3           0.109
```

BACK TO PRESENTATION

Ridge and elastic net models


```
ridge_logreg <- update(lasso_logreg, alpha = 0)
plot(ridge_logreg, xvar = "lambda")
```



```
pretty_coefs(coef(ridge_logreg, s = exp(2))) # all shrunk but no real ranking or anything
```

```
## # A tibble: 10 x 2
##   predictor coefficient
##   <chr>          <dbl>
## 1 (Intercept)    -0.630
## 2 V6              0.0457
## 3 V3              0.0448
## 4 V2              0.0447
## 5 V7              0.0414
## 6 V1              0.0395
## 7 V8              0.0391
## 8 V4              0.0383
## 9 V5              0.0373
## 10 V9             0.0224
```

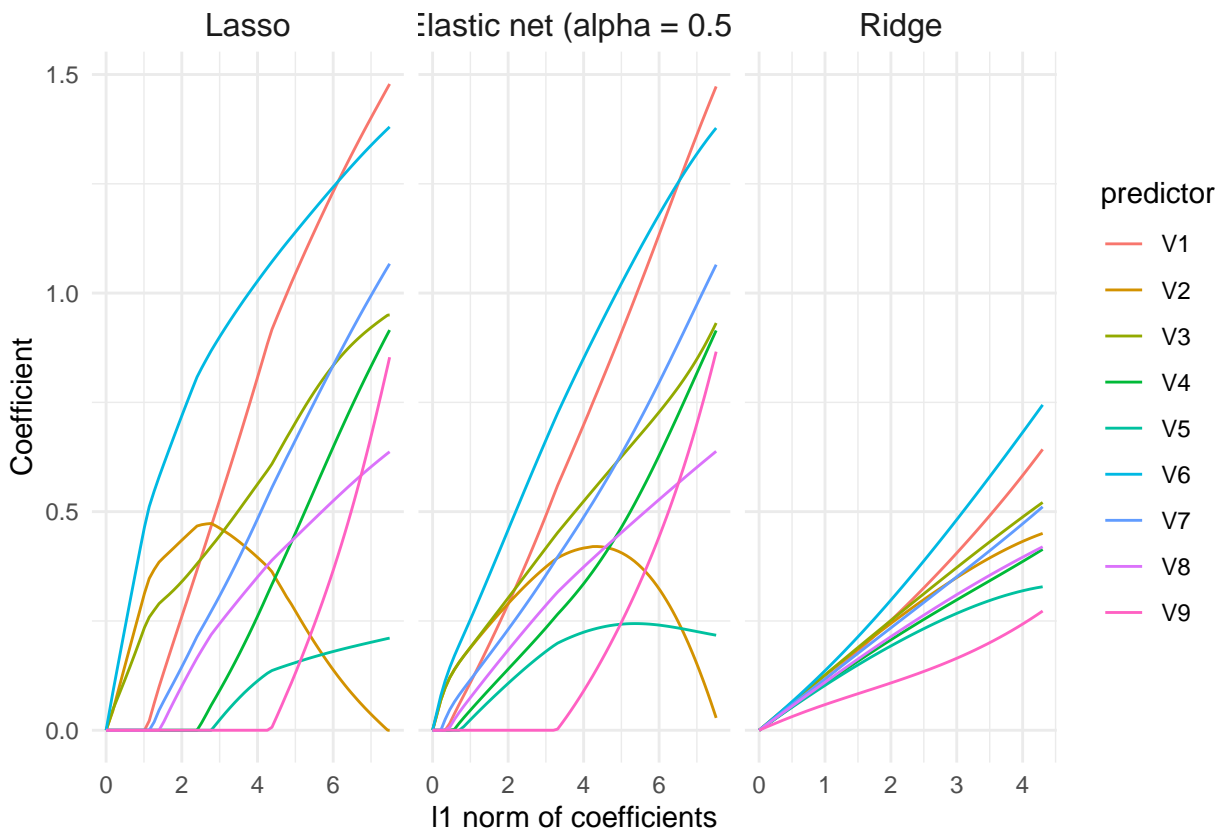
```
elastic_logreg <- update(lasso_logreg, alpha = 0.5)
plot(elastic_logreg, xvar = "lambda")
```



```
pretty_coefs(coef(elastic_logreg, s = exp(-1.2))) # all shrunk AND "ranking"
```

```
## # A tibble: 9 x 2
##   predictor coefficient
##   <chr>          <dbl>
## 1 (Intercept)   -0.719
## 2 V6             0.298
## 3 V3             0.215
## 4 V2             0.213
## 5 V1             0.148
## 6 V7             0.141
## 7 V8             0.0996
## 8 V4             0.0675
## 9 V5             0.0435
```

```
ldply(list(Lasso = lasso_logreg, Ridge = ridge_logreg, Elastic_net = elastic_logreg),
  function(.) with(., data.frame(as.matrix(t(beta)), x = apply(abs(beta), 2, sum))), .id = "mod") %>%
gather(predictor, value, -x, -mod) %>%
mutate(mod = factor(mod, levels = c("Lasso", "Elastic_net", "Ridge"),
  labels = c("Lasso", "Elastic net (alpha = 0.5)", "Ridge"))) %>%
ggplot(aes(x, value, colour = predictor)) +
  geom_line() +
  labs(x = "l1 norm of coefficients", y = "Coefficient") +
  facet_wrap(~ mod, scales = "free_x") +
  theme(strip.text = element_text(size = 12))
```

Over-fitting biopsy

It's a little data set but let's try and split into training and test sets. There is quite some over-fitting on the relative scale, but the absolute prediction error difference is minor.

```
set.seed(42) # reproducible stochastic code
train_idx <- runif(nrow(biopsy_complete)) <= 0.8 # not exactly indices
mean(train_idx) # fraction in training set

## [1] 0.806735

biopsy_train <- filter(biopsy_complete, train_idx)
biopsy_test <- filter(biopsy_complete, !train_idx)
all_equal(biopsy_complete, bind_rows(biopsy_train, biopsy_test)) # sanity check

## [1] TRUE

# Alternative with actual indices (mostly a matter of taste)
set.seed(42)
train_idx <- which(runif(nrow(biopsy_complete)) <= 0.8)
length(train_idx) / nrow(biopsy_complete)

## [1] 0.806735
```

```

biopsy_train <- slice(biopsy_complete, train_idx)
biopsy_test <- slice(biopsy_complete, -train_idx)
all_equal(biopsy_complete, bind_rows(biopsy_train, biopsy_test)) # sanity check

## [1] TRUE

# Normalise predictors and put them in matrix format
predictors_train <- select(biopsy_train, -ID, -class) %>%
  scale()
predictors_test <- select(biopsy_test, -ID, -class) %>%
  scale()

# Train model
lasso_biopsy_train <- glmnet(predictors_train, biopsy_train$class, family = "binomial")
D_train <- predict(lasso_biopsy_train, predictors_train, type = "class") %>%
  apply(2, function(.) mean(. != biopsy_train$class)) # s39 is the best-performing


D_test <- predict(lasso_biopsy_train, predictors_test, type = "class") %>%
  apply(2, function(.) mean(. != biopsy_test$class))

tibble(D_test = D_test[which.min(D_train)],
       D_train = D_train[which.min(D_train)],
       D_diff_abs = scales::percent(D_test - D_train),
       D_diff_rel = scales::percent((D_test - D_train) / D_train, big.mark = ","))

## # A tibble: 1 x 4
##   D_test D_train D_diff_abs D_diff_rel
##   <dbl>   <dbl> <chr>      <chr>
## 1 0.0379 0.0272 1%          39%


ggplot(mapping = aes(x = log(lasso_biopsy_train$lambda))) +
  coord_cartesian(ylim = c(0, NA)) +
  geom_line(aes(y = D_train, colour = "Within-sample")) +
  geom_line(aes(y = D_test, colour = "Hold-out")) +
  labs(y = "log(prediction error)", x = expression(log(lambda))) +
  scale_x_reverse()

```



Delassoing

```
plot(lasso_logreg, xvar = "lambda")
```



```
lambda <- exp(-4.5)
beta <- coef(lasso_logreg, x = biopsy_predictors, y = biopsy_complete$class,
             s = lambda/nrow(biopsy_complete), exact = TRUE)
delasso_fit <- fixedLassoInf(biopsy_predictors, as.numeric(biopsy_complete$class)-1, beta, lambda, "binomial",
                           alpha = 0.05)
```

```
## Warning in fixedLogitLassoInf(x, y, beta, lambda, alpha = alpha, type =
## type, : Solution beta does not satisfy the KKT conditions (to within specified
## tolerances)
```

```
delasso_fit
```


```
##
## Call:
## fixedLassoInf(x = biopsy_predictors, y = as.numeric(biopsy_complete$class) -
## 1, beta = beta, lambda = lambda, family = "binomial", alpha = 0.05)
##
## Testing results at lambda = 0.011, with alpha = 0.050
##
## Var    Coef Z-score P-value LowConfPt UpConfPt LowTailArea UpTailArea
## 1  1.509  3.770  0.596   -9.749   2.005    0.025    0.024
## 2 -0.019 -0.030  0.989    0.857    Inf     0.025    0.000
## 3  0.964  1.399  0.962   -Inf     0.570    0.000    0.025
## 4  0.947  2.680  0.808  -17.165   1.237    0.025    0.024
## 5  0.215  0.617  0.863  -10.079   0.635    0.025    0.024
## 6  1.396  4.084  0.000    0.743   3.114    0.025    0.025
## 7  1.095  2.611  0.749  -14.329   1.537    0.025    0.025
## 8  0.650  1.889  0.805  -12.461   1.008    0.025    0.025
## 9  0.927  1.629  0.705  -10.952   1.724    0.025    0.025
```

```
##
## Note: coefficients shown are full regression coefficients
```

Cross-validation

Use cross-validation to find best λ value. We found quite clear over-fitting above. Let's try to remedy this with cross-validation.

```
lasso_logreg_cv <- cv.glmnet(predictors_train, biopsy_train$class, family = "binomial")
plot(lasso_logreg_cv)
```



```
with(lasso_logreg_cv, data.frame(lambda.min, lambda.1se))


##      lambda.min lambda.1se
## 1 0.006531203 0.02893729

# With ggplot2 (more control and prettier)
cv_glmnet_coef_profile <- function(fit) {
  fade_text <- function(x, alpha = 0.2) {
    ifelse(paste(x) == lag(paste(x), default = ""), alpha, 1)
  }
  every_n <- function(x, n = 5) {
    seq_along(x) %% n == 0
  }
  with(fit, tibble(log_lambda = log(lambda), cvm, cvup, cvlo, nzero)) %>%
    ggplot(aes(x = log_lambda)) +
      geom_vline(xintercept = log(fit$lambda.min), linetype = 2, size = 0.5, colour = "red") +
```

```

    geom_vline(xintercept = log(fit$lambda.1se), linetype = 2, size = 0.5, colour = "dodgerblue") +
    geom_linerange(aes(ymin = cvlo, ymax = cvup), size = 0.3) +
    geom_ribbon(aes(ymin = cvlo, ymax = cvup), alpha = 0.1) +
    geom_point(aes(y = cvm), shape = 18, size = 1.5) +
    geom_text(aes(y = max(cvup) * 1.05, label = nzero, alpha = fade_text(nzero)), size = 8 / gg) +
    scale_alpha_identity() +
    labs(x = expression(log(lambda)), y = fit$name) +
    theme_minimal()
}
cv_glmnet_coef_profile(lasso_logreg_cv)

```



Test out-of-sample performance of the CV model

```

pred_min <- predict(lasso_logreg_cv, predictors_test, s = "lambda.min", type = "class")
pred_1se <- predict(lasso_logreg_cv, predictors_test, s = "lambda.1se", type = "class")

```

Cross-validation to pick best combination of λ and α

- grid search over alpha and lambda


Resampling methods: jackknife and bootstrap

Sampling distribution

```
age <- rnorm(1000, 60, 12)
age <- age + 2 * (age < 30)
```

```
age_samples <- ldply(setNames(1:400, 1:400), function(i) {
  set.seed(42 + i)
  enframe(sample(age, 40, replace = FALSE))
}, .id = "sample") %>%
  group_by(sample) %>%
  mutate(facet_title = sprintf("%s: Mean = %.1f", sample, mean(value)))
```

```
ggplot(filter(age_samples, sample %in% 1:20), aes(value)) +
  geom_histogram(bins = 10) +
  facet_wrap(~ facet_title, ncol = 5) +
  labs(x = "Age", y = "N")
```




```
aggregate_stats_age <- group_by(age_samples, sample) %>%
  summarise("Sample mean" = mean(value),
            "Sample variance" = var(value),
            "Sample 90-percentile" = quantile(value, 0.9))
```


```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
sample_plot <- llply(paste("Sample", c("mean", "variance", "90-percentile")),
  function(.) qplot(x = !!sym(.), data = aggregate_stats_age, geom = "histogram", bins = 10))


sample_plot[[1]]
```



```
sample_plot[[2]] + scale_x_continuous(sec.axis = sec_axis(sqrt, "Sample std.dev."))
```

```
sample_plot[[3]]
```



Jackknife

```
x <-c(8.26, 6.33, 10.4, 5.27, 5.35, 5.61, 6.12, 6.19, 5.2, 7.01, 8.74, 7.78, 7.02, 6, 6.5, 5.8, 5.12, 7)
coef_var <- function(x) sqrt(var(x))/mean(x) # coefficient of variation
coef_var(x)

## [1] 0.2524712

library(bootstrap)

##
## Attaching package: 'bootstrap'

## The following object is masked from 'package:broom':
##
## bootstrap

jackknife(x, coef_var)

## $jack.se
## [1] 0.05389943
##
## $jack.bias
## [1] -0.009266436
```

```
##
## $jack.values
## [1] 0.2563873 0.2565586 0.2384298 0.2507329 0.2513200 0.2530603 0.2557374
## [8] 0.2560293 0.2501992 0.2580969 0.2541045 0.2577524 0.2581067 0.2551946
## [15] 0.2571038 0.2541711 0.2495662 0.2581975 0.2571609 0.2561093 0.2020978
## [22] 0.2529980 0.2515338 0.2573745 0.2541045
##
## $call
## jackknife(x = x, theta = coef_var)
```

Exercise: cross-validation

```
library(MASS)
data(biopsy)
biopsy_complete <- na.exclude(biopsy)
summary(biopsy_complete)
```

```
##           ID              V1              V2              V3
## Length:683      Min.   : 1.000      Min.   : 1.000      Min.   : 1.000
## Class :character 1st Qu.: 2.000      1st Qu.: 1.000      1st Qu.: 1.000
## Mode  :character Median : 4.000      Median : 1.000      Median : 1.000
##              Mean   : 4.442      Mean   : 3.151      Mean   : 3.215
##              3rd Qu.: 6.000      3rd Qu.: 5.000      3rd Qu.: 5.000
##              Max.   :10.000      Max.   :10.000      Max.   :10.000
##           V4           V5           V6           V7
## Min.   : 1.00      Min.   : 1.000      Min.   : 1.000      Min.   : 1.000
## 1st Qu.: 1.00      1st Qu.: 2.000      1st Qu.: 1.000      1st Qu.: 2.000
## Median : 1.00      Median : 2.000      Median : 1.000      Median : 3.000
## Mean   : 2.83      Mean   : 3.234      Mean   : 3.545      Mean   : 3.445
## 3rd Qu.: 4.00      3rd Qu.: 4.000      3rd Qu.: 6.000      3rd Qu.: 5.000
## Max.   :10.00      Max.   :10.000      Max.   :10.000      Max.   :10.000
##           V8           V9           class
## Min.   : 1.00      Min.   : 1.000      benign   :444
## 1st Qu.: 1.00      1st Qu.: 1.000      malignant:239
## Median : 1.00      Median : 1.000
## Mean   : 2.87      Mean   : 1.603
## 3rd Qu.: 4.00      3rd Qu.: 1.000
## Max.   :10.00      Max.   :10.000
```

```
predictors <- biopsy_complete %>%
  select(-ID, -class)
pca_fit <- prcomp(predictors, scale = TRUE)
df_pca <- data.frame(pca_fit$x[, 1:4], outcome = biopsy_complete$class)
glm_fit <- glm(outcome ~ PC1 + PC2 + PC3 + PC4, data = df_pca, family = binomial)
summary(glm_fit)
```

```
##
## Call:
## glm(formula = outcome ~ PC1 + PC2 + PC3 + PC4, family = binomial,
##      data = df_pca)
##
## Deviance Residuals:
```

```
##      Min      1Q   Median      3Q      Max
## -3.1791 -0.1304 -0.0619  0.0228  2.4799
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -1.0739     0.3035  -3.539 0.000402 ***
## PC1          -2.4140     0.2556  -9.445 < 2e-16 ***
## PC2          -0.1592     0.5050  -0.315 0.752540
## PC3           0.7191     0.3273   2.197 0.028032 *
## PC4          -0.9151     0.3691  -2.479 0.013159 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 884.35  on 682  degrees of freedom
## Residual deviance: 106.12  on 678  degrees of freedom
## AIC: 116.12
##
## Number of Fisher Scoring iterations: 8
```

```
tidy(glm_fit)
```

```
## # A tibble: 5 x 5
##   term      estimate std.error statistic  p.value
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept) -1.07      0.303    -3.54 4.02e- 4
## 2 PC1        -2.41      0.256    -9.44 3.56e-21
## 3 PC2        -0.159    0.505    -0.315 7.53e- 1
## 4 PC3         0.719    0.327     2.20 2.80e- 2
## 5 PC4        -0.915    0.369    -2.48 1.32e- 2
```

1. LOO-CV error rate

```
library(boot)
glm_fit_loocv <- cv.glm(df_pca, glm_fit)
glm_fit_loocv$delta
```

```
## [1] 0.02301890 0.02301788
```

2. Use proper cost function

```
glm_fit_loocv2 <- cv.glm(df_pca, glm_fit, cost = function(r, pi = 0) mean(abs(r-pi) > 0.5))
glm_fit_loocv2$delta
```

```
## [1] 0.02781845 0.02785918
```


3. Difference between error rates, and their interpretation

```
# The binary cost function forces the predictions into a binary class assignment (w.r.t. the 0.5 thresh
```

```

expand.grid(y_obs = 0:1,
            y_pred = 0:100 / 100) %>%
  rowwise() %>%
  mutate(cost_mse= mean((y_pred - y_obs)^2),
         cost_binary = mean(abs(y_pred - y_obs) > 0.5),
         cost_mae = mean(abs(y_pred - y_obs))) %>%
  pivot_longer(starts_with("cost_"), names_to = "cost_fun", values_to = "cost") %>%
  ggplot(aes(y_pred, cost, colour = cost_fun)) +
    geom_line() +
    facet_wrap(~ y_obs)

```



4. 10-fold CV

```

# The error rates change quite a bit, which makes sense because 10-fold CV has a lot fewer folds than L
glm_fit_10cv <- update(glm_fit_loocv, K = 10)
glm_fit_10cv$delta

```

```
## [1] 0.02412049 0.02399288
```

```

# glm_fit_10cv2 <- cv.glm(df_pca, glm_fit, cost = function(r, pi = 0) mean(abs(r-pi) > 0.5), K = 10)
glm_fit_10cv2 <- update(glm_fit_loocv2, K = 10)
glm_fit_10cv2$delta

```

```
## [1] 0.03074671 0.03118830
```

```
ldply(list("L00" = glm_fit_loocv, "L00 2" = glm_fit_loocv2, "10-fold" = glm_fit_10cv, "10-fold 2" = glm_fit_10cv2,
  with, delta) %>%
  setNames(c("model", "error_rate", "corrected_error_rate"))

##      model error_rate corrected_error_rate
## 1      L00 0.02301890          0.02301788
## 2      L00 2 0.02781845          0.02785918
## 3 10-fold 0.02412049          0.02399288
## 4 10-fold 2 0.03074671          0.03118830

knitr::opts_chunk$set(include = FALSE)
```

Exercise: penalised regression

1. + 2. Lasso regression
3. Why does it normally make sense to normalise predictors
4. Using CV to get reasonable estimate of λ
5. Obtain coefficients for “best” λ
6. Re-fit with correct family

Quite some difference between the sets of predictors kept:

7. As ridge regression
8. Get idea about sparse solution using ridge results?
9. Elastic net ($\alpha = 0.5$)

Ideally, one should do CV over α as well.

10. DeLasso the results
11. Selective inference